

```
Descriptions_tools = [
{
  "name": "annotate_open_reading_frames",
  "description": (
    "Find all Open Reading Frames (ORFs) in a DNA sequence using Biopython."
    "Scans the forward strand and (optionally) the reverse-complement strand,"
    "returning ORF coordinates that meet a minimum nucleotide length."
  ),
  "required_parameters": [
    {
      "name": "sequence",
      "type": "str",
      "default": None,
      "description": (
        "DNA sequence to analyze. Should be provided as a raw string of nucleotides"
        "(typically A/C/G/T; clarify whether ambiguous bases like N are allowed)."
      ),
    },
    {
      "name": "min_length",
      "type": "int",
      "default": None,
      "description": (
        "Minimum ORF length in nucleotides (not amino acids). ORFs shorter than this are excluded."
      )
    }
  ]
}
```

```
        ),  
    },  
],  
"optional_parameters": [  
    {  
        "name": "search_reverse",  
        "type": "bool",  
        "default": False,  
        "description": (  
            "If True, also scan the reverse-complement strand. This typically increases runtime "  
            "by ~2x on the same input length."  
        ),  
    },  
    {  
        "name": "filter_subsets",  
        "type": "bool",  
        "default": False,  
        "description": (  
            "If True, filter out ORFs that share the same end position but have a later start "  
            "(i.e., remove subset ORFs nested at the same stop site)."  
        ),  
    },  
],  
"hardware_requirements": {  
    "device": "cpu_only",  
    "notes": "No GPU required. Performance depends mainly on sequence length and CPU single-thread speed.",
```

```
},
"time_complexity": {
    "assumptions": (
        "Report wall-clock latency for a typical input size (e.g., 5–10 kb plasmid, 1 Mb contig, etc.),"
        "
        "and specify CPU model + threads. Include cold-start vs warm-start if Biopython import
        dominates."
    ),
    "latency_seconds": {"n1": None, "n2": None, "n10": None}, # doc-writer fills with
    benchmarks
},
{
```

```
{
    "name": "annotate_plasmid",
    "description": (
        "Annotate a DNA sequence by invoking pLannotate via its command-line interface. "
        "Typically used for plasmid annotation (features, CDS, common parts) and supports "
        "circular sequences."
    ),
    "required_parameters": [
        {
            "name": "sequence",
            "type": "str",
            "default": None,
            "description": (
                "DNA sequence to annotate (raw nucleotide string). Clarify expected length range "
                "and whether FASTA headers/newlines are accepted."
            )
        }
    ]
}
```

```
        ),
    }
],
"optional_parameters": [
{
    "name": "is_circular",
    "type": "bool",
    "default": True,
    "description": (
        "If True, treat the sequence as circular during annotation (important for plasmids). "
        "If False, treat as linear DNA."
    ),
}
],
"hardware_requirements": {
    "device": "cpu_only",
    "notes": (
        "No GPU required. Requires pLannotate installed and accessible in PATH. "
        "May need substantial RAM depending on database/index usage."
    ),
}
],
"time_complexity": {
    "assumptions": (
        "Benchmark includes CLI startup + any database/index loading. Provide times for typical
plasmid sizes "
        "(e.g., 3–20 kb) and specify CPU + storage type (network FS vs local SSD can matter)."
    ),
}
```

```
    "latency_seconds": {"n1": None, "n2": None, "n10": None}, # doc-writer fills with
    benchmarks
  },
}
```

Name

What it is:

The unique function/tool identifier.

Complete it by:

- Matching the exact function or CLI wrapper name in code.
- Keeping it stable across versions (avoid renaming).
- Using snake_case and being explicit (e.g., annotate_open_reading_frames, not orf_tool).

Why it matters:

Agents and pipelines call this string directly — any mismatch breaks execution.

description

What it is:

A high-level explanation of what the tool does, suggestions on when to use to tool and explanation about how good the tool is compared to the others

Complete it by:

- Stating input → processing → output clearly.
- Naming the underlying library/algorithm (Biopython, CLI, model, etc.).
- Mentioning key behavior assumptions (strand search, circular DNA, etc.).
- Keeping it concise (2–4 sentences max).

Why it matters:

Helps humans and agents choose the correct tool without reading code.

required_parameters

What it is:

Inputs that must always be provided by the agent for the tool.

Complete it by (for each parameter):

- name → exact argument name in code
- type → strict type (str, int, bool, etc.)
- description → what it means + units + format
- default → usually None, but if a value is said by the GitHub to be fixed set the default value

Also specify:

- expected format (FASTA vs raw string)
- units (nt vs aa)
- allowed ranges
- constraints or validation rules

Why it matters:

Prevents silent failures and incorrect inputs.

optional_parameters

What it is:

Inputs that modify behavior but aren't required.

Complete it by (for each parameter):

- real default value (must match code)
- clear explanation of what changes when enabled/disabled
- note performance or memory impact if applicable

Why it matters:

Prevents ambiguous behavior and hidden runtime costs.

Hardware_requirements

What it is:

The compute environment needed to run the tool.

Complete it by specifying:

- cpu_only / gpu_optional / gpu_required

Example:

- CPU only

Why it matters:

Lets schedulers, clusters, or agents route jobs correctly.

Time_complexity

What it is:

Real-world runtime measurements.

Complete it by:

- Measuring wall-clock time (not Big-O math)
- Stating assumptions:
 - hardware used
 - input size
 - threads/batching
- reporting:
 - time for 1 run
 - time for 2 runs
 - time for 10 runs

Example:

- 1 → 0.4s
- 2 → 0.7s
- 10 → 3.2s

Why it matters:

Agents can plan batching, parallelization, and scheduling accurately.

Outputs

What it is:

What the tool returns.

Complete it by specifying:

- data type (list, dict, dataframe, file path, JSON)
- schema/fields
- units/format
- example output

Why it matters:

Downstream tools must parse this reliably.

Failure_modes

What it is:

Common errors and how to handle them.

Complete it by listing:

- error condition
- likely cause
- recommended fix

Example:

- empty sequence → invalid input → validate before call

Why it matters:

Prevents silent crashes in pipelines.

We also will require requirements for each tool and function call:

Dependencies should be in a separate .txt file (requirements_Function_name.txt)

What it is:

External libraries, binaries, or models required.

Complete it by specifying:

- package names + versions
- installation method

- whether system binary or Python package

Why it matters:

Ensures reproducibility across machines/HPC nodes.